

Dhvani Dave

NY | +1 716-553-7650 | dhvanidave7014@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Software Engineer who enjoys building meaningful products, solving problems and learning along the way.

EDUCATION

University at Buffalo

Master of Science in Computer Science; GPA: 3.9/4.0

Buffalo, NY

Aug. 2025 – Dec. 2026

University of Mumbai

Bachelor of Engineering in Computer Engineering; GPA: 3.6/4.0

Mumbai, India

Aug. 2019 – Jun. 2023

TECHNICAL SKILLS

Languages: Python, Go, Java, C/C++, TypeScript, JavaScript, SQL

AI/ML: Machine Learning, Deep Learning, PyTorch, TensorFlow, Hugging Face, NumPy, Pandas

LLM & RAG: LangGraph, RAG, FAISS, Vector Search, Embeddings, LLM Fine-Tuning (QDoRA)

Backend & APIs: FastAPI, Node.js, REST, GraphQL, Microservices, Concurrent Programming

Frontend: React.js, Next.js, Streamlit

Data & Infrastructure: PostgreSQL, MongoDB, Neo4j, AWS (EC2, S3, Lambda), Docker, Kubernetes, Git, GitHub Actions, CI/CD, Linux

EXPERIENCE

Cloud Software Engineer

Jul. 2023 – Apr. 2025

CloudRock

Mumbai, India

- Engineered and deployed a full enterprise **Workday** system implementation for a major insurance client, owning the technical lifecycle from requirements analysis through configuration, integration, and **go-live**.
- Built and maintained Workday integrations using **Workday Studio, XSLT, and XML**, developing custom **EIB templates** to automate data exchange between Workday and downstream HR/Finance systems.
- Engineered custom **calculated fields and business process rules** to streamline integration logic, reducing data errors by **30%** and cutting integration processing time.
- Designed and executed **unit, integration, and User Acceptance Testing (UAT)** suites, debugging configuration logic early to catch defects and reduce downstream rework before production.
- Shipped the system to production **on time and within budget with zero production defects**, reflecting rigorous testing discipline and strong attention to system reliability and scale.
- Partnered directly with client stakeholders to gather requirements and validate system behavior, ensuring engineering decisions stayed aligned with real business workflows and long-term scalability.
- Owned **post-launch support** and technical troubleshooting, diagnosing and resolving functional bugs to ensure a smooth, stable transition to production for end users.
- Authored internal engineering documentation and led **knowledge-sharing sessions** to mentor teammates, streamlining onboarding and improving delivery consistency across future projects.

PROJECTS

CiteMind – Citation-Grounded Research Paper Assistant | FastAPI, RAG, LLM, Vector Search

Jul. 2026

- Built a **FastAPI-based RAG pipeline** for processing research papers, including **PDF extraction, text chunking, embedding generation**, retrieval, and grounded answer generation.
- Implemented **semantic chunking and vector embedding search** to ground LLM responses in retrieved passages, improving accuracy and reducing hallucinations with **traceable, source-cited answers**.

SentinelAI – Coral Hackathon | FastAPI, Next.js, GraphQL, PostgreSQL, Neo4j

May 2026

- Architected a **distributed 6-agent backend pipeline** (FastAPI) integrating **PostgreSQL** and **Neo4j** for **graph-based fraud detection**, exposing a **GraphQL API** consumed by the Next.js frontend for investigation reporting.
- Validated the system end-to-end on a case with **8 high-risk transactions**, producing a **fraud risk score of 100** with accurate relationship detection across the transaction graph.

CodeSense – Agentic AI Developer Copilot | FastAPI, React, LangGraph, FAISS, Tree-sitter, Qwen3-8B

Mar. 2026

- Designed, coded, and deployed a **full-stack agentic debugging assistant** (FastAPI backend, React frontend) that retrieves code context, detects bugs, and generates validated patches across the full development lifecycle.
- Experimented with large language models (LLMs) by **fine-tuning Qwen3-8B with QDoRA** on **5,400 samples**, improving **SWE-bench Token F1 from 0.075 to 0.257** and **reducing hallucinations by 93%**.

PROFESSIONAL CERTIFICATIONS

BCG GenAI Job Simulation | Forage

Aug. 2026

- Built an **AI-powered financial chatbot** for BCG's GenAI Consulting team using Python and pandas.

Build RAG Applications: Get Started | IBM

Aug. 2026

- Covered core RAG concepts, including embeddings, vector search, and integrating retrieved context into LLM-powered applications.

Introduction to Model Context Protocol (MCP) | Anthropic

Aug. 2026

- Studied **MCP** for standardized communication between AI models, tools, and external data sources, including **client/server architecture** and secure context sharing.

Building with the Claude API | Anthropic

Jul. 2026

- Hands-on training building AI-powered applications, handling **API requests, structured outputs**, and prompt workflows.